

```
!pip install -q tensorflow-datasets tensorflow matplotlib

import numpy as np
import tensorflow as tf

import tensorflow_datasets as tfds

from tensorflow.keras.utils import to_categorical

## Loading images and labels
(train_ds, train_labels), (test_ds, test_labels) = tfds.load("tf_flowers",
    split=["train[:70%]", "train[:30%]"], ## Train test split
    batch_size=-1,
    as_supervised=True, # Include labels
)

Downloading and preparing dataset 218.21 MiB (download: 218.21 MiB, generated: 221.83 MiB, total: 440.05 MiB) to /root/tensorflow_datasets/tf_flowers/3.0.1...
Dl Completed...: 100% 5/5 [00:02<00:00, 2.01 file/s]
Dataset tf_flowers downloaded and prepared to /root/tensorflow_datasets/tf_flowers/3.0.1. Subsequent calls will reuse this data.

## check existing image size
train_ds[0].shape

TensorShape([442, 1024, 3])

## Resizing images
train_ds = tf.image.resize(train_ds, (150, 150))
test_ds = tf.image.resize(test_ds, (150, 150))

train_labels

<tf.Tensor: shape=(2569,), dtype=int64, numpy=array([2, 3, 3, ..., 0, 2, 0])>

## Transforming labels to correct format
train_labels = to_categorical(train_labels, num_classes=5)
test_labels = to_categorical(test_labels, num_classes=5)

from tensorflow.keras.applications.vgg16 import VGG16
from tensorflow.keras.applications.vgg16 import preprocess_input

train_ds[0].shape

TensorShape([150, 150, 3])

## Loading VGG16 model
base_model = VGG16(weights="imagenet", include_top=False, input_shape=train_ds[0].shape)

Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/vgg16/vgg16_weights_tf_dim_ordering_tf_kernels_notop.h5
58889256/58889256 [=====] - 0s 0us/step

## will not train base mode
# Freeze Parameters in model's lower convolutional layers
base_model.trainable = False

## will not train base mode
# Freeze Parameters in model's lower convolutional layers
base_model.trainable = False

## Preprocessing input
train_ds = preprocess_input(train_ds)
test_ds = preprocess_input(test_ds)

## model details
base_model.summary()
```

Model: "vgg16"

Layer (type)	Output Shape	Param #

input_1 (InputLayer)	[(None, 150, 150, 3)]	0
block1_conv1 (Conv2D)	(None, 150, 150, 64)	1792
block1_conv2 (Conv2D)	(None, 150, 150, 64)	36928
block1_pool (MaxPooling2D)	(None, 75, 75, 64)	0
block2_conv1 (Conv2D)	(None, 75, 75, 128)	73856
block2_conv2 (Conv2D)	(None, 75, 75, 128)	147584
block2_pool (MaxPooling2D)	(None, 37, 37, 128)	0
block3_conv1 (Conv2D)	(None, 37, 37, 256)	295168

```

block3_conv2 (Conv2D)      (None, 37, 37, 256)      590080
block3_conv3 (Conv2D)      (None, 37, 37, 256)      590080
block3_pool (MaxPooling2D) (None, 18, 18, 256)      0
block4_conv1 (Conv2D)      (None, 18, 18, 512)      1180160
block4_conv2 (Conv2D)      (None, 18, 18, 512)      2359808
block4_conv3 (Conv2D)      (None, 18, 18, 512)      2359808
block4_pool (MaxPooling2D) (None, 9, 9, 512)        0
block5_conv1 (Conv2D)      (None, 9, 9, 512)        2359808
block5_conv2 (Conv2D)      (None, 9, 9, 512)        2359808
block5_conv3 (Conv2D)      (None, 9, 9, 512)        2359808
block5_pool (MaxPooling2D) (None, 4, 4, 512)        0

```

```

=====
Total params: 14714688 (56.13 MB)
Trainable params: 0 (0.00 Byte)
Non-trainable params: 14714688 (56.13 MB)

```

```

#add our layers on top of this model
from tensorflow.keras import layers, models

```

```

flatten_layer = layers.Flatten()
dense_layer_1 = layers.Dense(50, activation='relu')
dense_layer_2 = layers.Dense(20, activation='relu')
prediction_layer = layers.Dense(5, activation='softmax')

```

```

model = models.Sequential([
    base_model,
    flatten_layer,
    dense_layer_1,
    dense_layer_2,
    prediction_layer
])

```

```

from tensorflow.keras.callbacks import EarlyStopping

```

```

model.compile(
    optimizer='adam',
    loss='categorical_crossentropy',
    metrics=['accuracy'],
)

```

```

es = EarlyStopping(monitor='val_accuracy', mode='max', patience=5, restore_best_weights=True)

```

```

history=model.fit(train_ds, train_labels, epochs=5, validation_split=0.2, batch_size=32, callbacks=[es])

```

```

Epoch 1/5
65/65 [=====] - 657s 10s/step - loss: 0.8456 - accuracy: 0.7071 - val_loss: 1.2420 - val_accuracy: 0.6304
Epoch 2/5
65/65 [=====] - 647s 10s/step - loss: 0.5961 - accuracy: 0.7898 - val_loss: 1.0110 - val_accuracy: 0.6790
Epoch 3/5
65/65 [=====] - 630s 10s/step - loss: 0.3979 - accuracy: 0.8589 - val_loss: 1.1255 - val_accuracy: 0.6984
Epoch 4/5
65/65 [=====] - 630s 10s/step - loss: 0.3205 - accuracy: 0.8818 - val_loss: 1.1503 - val_accuracy: 0.7101
Epoch 5/5
65/65 [=====] - 631s 10s/step - loss: 0.2371 - accuracy: 0.9119 - val_loss: 1.3278 - val_accuracy: 0.6946

```

```

los,accurac=model.evaluate(test_ds,test_labels)
print("Loss: ",los,"Accuracy: ", accurac)

```

```


35/35 [=====] - 291s 8s/step - loss: 0.1887 - accuracy: 0.9292
Loss: 0.18868598341941833 Accuracy: 0.9291552901268005

```

```

import matplotlib.pyplot as plt
plt.plot(history.history['accuracy'])
plt.title('ACCURACY')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train'],loc='upper left')
plt.show()

```



```
import numpy as np
import pandas as pd
y_pred = model.predict(test_ds)
y_classes = [np.argmax(element) for element in y_pred]
#to_categorical(y_classes, num_classes=5)
#to_categorical(test_labels, num_classes=5)
print(y_classes[:10])
print("\nTest")
print(test_labels[:10])
```

```
35/35 [=====] - 262s 7s/step
[2, 3, 3, 2, 3, 0, 0, 0, 0, 4]
```

```
Test
[[0. 0. 1. 0. 0.]
 [0. 0. 0. 1. 0.]
 [0. 0. 0. 1. 0.]
 [0. 0. 0. 0. 1.]
 [0. 0. 0. 1. 0.]
 [1. 0. 0. 0. 0.]
 [1. 0. 0. 0. 0.]
 [1. 0. 0. 0. 0.]
 [1. 0. 0. 0. 0.]
 [0. 1. 0. 0. 0.]]
```